

Model Monitoring Pipeline and Drift Detection for ASR Microservice on AWS

Deploying the `wav2vec2-large-960h` Automatic Speech Recognition (ASR) model as a containerized microservice in the cloud introduces the need for robust monitoring to ensure reliability, detect performance degradation, and identify model drift over time. Given the project's EC2- and Docker-based architecture, a custom monitoring pipeline on Amazon Web Services (AWS) can be built using native AWS tools and open-source libraries to address these requirements without relying on SageMaker Model Monitor.

Monitoring Objectives

The monitoring pipeline must address three core concerns:

- **Operational health:** API uptime, latency, and resource usage.
- **Model performance:** Quality and reliability of transcriptions.
- **Model drift:** Detecting changes in input data or model behavior over time.

Architecture Overview

The monitoring pipeline integrates with the existing system:

- ASR microservice containerized and hosted on **AWS EC2**,
- **Elasticsearch** backend,
- **Search UI** as the end-user interface,
- Audio data stored in **Amazon S3** for auditability and retraining.

Monitoring with Cloudwatch

Each EC2 instance will be configured with the **CloudWatch Agent** to push logs and metrics, including API response times, status codes, memory and CPU usage.

CloudWatch Dashboards visualize these metrics, while **CloudWatch Alarms** notify engineers via **Amazon SNS** of anomalies (e.g., spike in failures, increased latency).

Custom Model Drift Detection

Given that `wav2vec2` processes unstructured audio, drift monitoring requires custom analysis.

a. Data Drift Detection

Audio files received by the ASR API will be preprocessed using `librosa` or `pydub` to extract audio-level features such as Duration, loudness (RMS), silence ratio.

These are stored in **Amazon Timestream** or **S3** for temporal analysis. Weekly or daily jobs (triggered by **AWS Lambda**) compare current distributions with historical baselines using statistical tests to detect significant deviations.

b. Concept Drift Detection

Where ground-truth transcriptions are available, **Word Error Rate (WER)** and **Levenshtein distance** are calculated. For production data, drift is inferred using:

- Transcription entropy,
- Out-of-vocabulary token frequency,
- Change in common n-gram patterns.

These metrics are logged and visualized to flag semantic inconsistencies.

Audit Logging and Feedback Loop

All transcriptions, audio metadata, and extracted features are stored in **Amazon S3**, **Athena** or **Glue** can be used to query this data. This supports both manual review and retraining pipelines.

Model Retraining

If drift or degradation is confirmed, a retraining pipeline can be initiated manually or scheduled. Audio samples flagged for drift or poor performance are re-labeled and added to an updated dataset for finetuning wav2vec2, with retraining occurring either locally or later migrated to **SageMaker Training Jobs**.

Conclusion

A well-integrated AWS-native monitoring pipeline, using EC2, CloudWatch, S3, and open-source audio libraries, enables proactive drift detection and robust maintenance of ASR service quality. This custom approach aligns with the containerized, low-overhead architecture while ensuring scalability, auditability, and long-term performance.